

Strings += Performance

Ryan Dougherty

Quick Refresher

- Question: Are Strings in Java immutable?
- Answer: Yes!
- More on this later...

Background

- Junior in Computer Science + Math minor
- Honors Thesis
- Extension of previous thesis

Honors Thesis

- Each file: $448 * 304$ cells, 4 bytes/cell
- 27 years * 52 weeks * 545 KB ~ 750 MB data
- 136k vertices $\rightarrow \sim 10^9$ (billion) edges
- $\sim 10^{12}$ (trillion) correlation computations
- Goal: make code run faster!

Original Code

```
byte[] byteArray = new byte[4];
try {
    File file = new File("binary_input_file");
    FileInputStream file_input = new FileInputStream(file);
    DataInputStream data_in = new DataInputStream(file_input);

    while (true) {
        try {
            for (int index = 0; index < 4; index++) {
                byteArray[index] = data_in.readByte();
            }
        } catch (EOFException eof) {
            break;
        }

        // transform byteArray to a float value
        float f = readFloatLittleEndian(byteArray);

        //... Do stuff with f...
    }
    data_in.close();
} catch (IOException e) {
    System.err.println(e.toString());
}
```


Execution Time?

- ~0.43 sec
- Not that fast, but a good baseline.

1 Code Change

- Floats easier than bytes
- Idea: write out all floats to file...How?
- Solution: Use a concatenated String!
- What was added in loop: `s += f + "\n";` // s is a String, f is a float

Execution Time?

- ~72 sec
- ~170x slower!

Why?

- Question: $1 + 2 + 3 + \dots + n = ?$
- $= n*(n+1)/2 = n^2/2 + n/2$
- $\in O(n^2)$

Why? (cont.)

- Reminder: `s += f + "\n";` // s is a String, f is a float
- Reason 1: Immutable Strings!
- Reason 2: Compiler cannot guess!
- Each assignment creates a new String of 1 character, then 2, etc.
- $O(n)$ loop $\rightarrow O(n^2)$ loop

Solution...?

- `java.lang.StringBuilder!` (and its `.append()` method)
- S.B. is interesting, ask afterward...
- Execution time: ~0.52-0.55 sec → ~20-28% slower
- Much better than 17000% slower!

Summary

	Baseline	String concatenation	StringBuilder
Execution Time	0.43 s	~72 s	0.52 s
Memory Usage (normalized)	0 MB	+155.68 MB	+3.86 MB

Take-home Messages

- Your code does not live in a vacuum!
- View the Java docs! (for fun...?)
- Analyze your code!
- Use profilers!

Suggestions

- View (compiled) object code: “javap -c MyClass”
- Use S.B. (not String) for “unknowable” data values.

Questions?

